

lab0-pandas-auto_mpg

September 26, 2022

1 Pandas

As described at <https://pandas.pydata.org> > pandas is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.

1.1 Resources

1. Ch 5-6 in Python for Data Analysis, 2nd Ed, Wes McKinney (UCalgary library and <https://github.com/wesm/pydata-book>)
2. Ch 3 in Python Data Science Handbook, Jake VanderPlas (Ucalgary library and <https://github.com/jakevdp/PythonDataScienceHandbook>)

Let's explore some of the features.

First, import Pandas, and Numpy as a good companion.

```
[1]: import numpy as np
import pandas as pd
```

1.2 Create pandas DataFrames

There are several ways to create Pandas DataFrames, most notably from reading a csv (comma separated values file). DataFrames are 'spreadsheets' in Python. We will often use `df` as a variable name for a DataFrame.

If data is not stored in a file, a DataFrame can be created from a dictionary of lists

```
data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],
        'year': [2000, 2001, 2002, 2001, 2002, 2003],
        'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}
frame = pd.DataFrame(data)
```

where dictionary keys become column headers.

An alternative is to create from a numpy array and set column headers separately:

```
[2]: # From a numpy array
df = pd.DataFrame( np.arange(20).reshape(5,4), columns=['alpha', 'beta', 'gamma', 'delta'])
```

```
df
```

```
[2]:
```

	alpha	beta	gamma	delta
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19

```
[3]: # checking its type
type(df)
```

```
[3]: pandas.core.frame.DataFrame
```

1.3 Indexing

Accessing data in Dataframes is done by rows and columns, either index or label based.

```
[4]: # select a column
df['alpha']
```

```
[4]:
```

0	0
1	4
2	8
3	12
4	16

Name: alpha, dtype: int64

```
[5]: # select two columns
df[['alpha', 'gamma']]
```

```
[5]:
```

	alpha	gamma
0	0	2
1	4	6
2	8	10
3	12	14
4	16	18

```
[6]: # select rows
df.iloc[:2]
```

```
[6]:
```

	alpha	beta	gamma	delta
0	0	1	2	3
1	4	5	6	7

```
[7]: # select rows and columns
df.iloc[:2, :2]
```

```
[7]:   alpha  beta
     0     0     1
     1     4     5
```

```
[8]: # select rows and columns, mixed
     df.loc[:2, ['alpha', 'beta']]
```

```
[8]:   alpha  beta
     0     0     1
     1     4     5
     2     8     9
```

1.4 DataFrame math

Similar to Numpy, DataFrames support direct math

```
[9]: # direct math
     df2 = (9/5) * df + 32
     df2
```

```
[9]:   alpha  beta  gamma  delta
     0  32.0  33.8   35.6   37.4
     1  39.2  41.0   42.8   44.6
     2  46.4  48.2   50.0   51.8
     3  53.6  55.4   57.2   59.0
     4  60.8  62.6   64.4   66.2
```

```
[10]: # add two dataframes of same shape
     df + df2
```

```
[10]:   alpha  beta  gamma  delta
     0  32.0  34.8   37.6   40.4
     1  43.2  46.0   48.8   51.6
     2  54.4  57.2   60.0   62.8
     3  65.6  68.4   71.2   74.0
     4  76.8  79.6   82.4   85.2
```

```
[11]: # map a function to each column
     f = lambda x: x.max() - x.min()

     df.apply(f)
```

```
[11]: alpha    16
     beta    16
     gamma    16
     delta    16
     dtype: int64
```

1.5 DataFrame manipulation

Adding and deleting columns, as well as changing entries is similar to Python dictionaries.

Note that most DataFrame methods do not change the DataFrame directly, but return a new DataFrame. It is always good to check how the method you are invoking behaves.

```
[12]: # add a column
df['epsilon'] = ['low', 'medium', 'low', 'high', 'high']
df
```

```
[12]:
```

	alpha	beta	gamma	delta	epsilon
0	0	1	2	3	low
1	4	5	6	7	medium
2	8	9	10	11	low
3	12	13	14	15	high
4	16	17	18	19	high

```
[13]: # What is the size?
df.shape
```

```
[13]: (5, 5)
```

```
[14]: # delete column
df_dropped = df.drop(columns=['gamma'])
df_dropped
```

```
[14]:
```

	alpha	beta	delta	epsilon
0	0	1	3	low
1	4	5	7	medium
2	8	9	11	low
3	12	13	15	high
4	16	17	19	high

```
[15]: # the original dataframe is unaffected
df
```

```
[15]:
```

	alpha	beta	gamma	delta	epsilon
0	0	1	2	3	low
1	4	5	6	7	medium
2	8	9	10	11	low
3	12	13	14	15	high
4	16	17	18	19	high

Let's create a copy and assign new values to the first column:

```
[16]: df_copy = df.copy()
df_copy['alpha'] = 20
print(df)
```

```
print(df_copy)
```

	alpha	beta	gamma	delta	epsilon
0	0	1	2	3	low
1	4	5	6	7	medium
2	8	9	10	11	low
3	12	13	14	15	high
4	16	17	18	19	high

	alpha	beta	gamma	delta	epsilon
0	20	1	2	3	low
1	20	5	6	7	medium
2	20	9	10	11	low
3	20	13	14	15	high
4	20	17	18	19	high

DataFrames can be sorted by column:

```
[17]: # sorting values
df.sort_values(by='epsilon')
```

```
[17]:
```

	alpha	beta	gamma	delta	epsilon
3	12	13	14	15	high
4	16	17	18	19	high
0	0	1	2	3	low
2	8	9	10	11	low
1	4	5	6	7	medium

1.6 Load data from file

Most often data will come from somewhere, often csv files, and using `pd.read_csv()` will allow smooth creation of DataFrames.

Let's load that same heart-attack.csv that we used in Numpy before:

```
[18]: data = pd.read_fwf('auto-mpg.data',
    ↳ names=['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
    ↳ 'acceleration', 'model year', 'origin', 'car name'])
```

After loading data, it is good practice to check what we have. Usually, the sequences is: 1. Check dimension 2. Peek at the first rows 3. Get info on data types and missing values 4. Summarize columns

```
[19]: # Check dimension (rows, columns)
data.shape
```

```
[19]: (398, 9)
```

```
[20]: # Peek at the first rows
data.head()
```

```
[20]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0   18.0          8         307.0        130.0   3504.0          12.0         70
1   15.0          8         350.0        165.0   3693.0          11.5         70
2   18.0          8         318.0        150.0   3436.0          11.0         70
3   16.0          8         304.0        150.0   3433.0          12.0         70
4   17.0          8         302.0        140.0   3449.0          10.5         70

      origin          car name
0         1  "chevrolet chevelle malibu"
1         1      "buick skylark 320"
2         1  "plymouth satellite"
3         1      "amc rebel sst"
4         1      "ford torino"
```

```
[21]: # Column names are
data.columns
```

```
[21]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
          'acceleration', 'model year', 'origin', 'car name'],
          dtype='object')
```

```
[22]: # Get info on data types and missing values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       398 non-null   object
4   weight           398 non-null   float64
5   acceleration     398 non-null   float64
6   model year      398 non-null   int64
7   origin           398 non-null   int64
8   car name         398 non-null   object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

1.7 Summarize values

What is the mean, std, min, max in each column?

```
[23]: data.mean()
```

```
/tmp/nix-shell.G00woP/ipykernel_46966/531903386.py:1: FutureWarning: Dropping of
nuisance columns in DataFrame reductions (with 'numeric_only=None') is
deprecated; in a future version this will raise TypeError.  Select only valid
columns before calling the reduction.
```

```
data.mean()
```

```
[23]: mpg                23.514573
      cylinders          5.454774
      displacement      193.425879
      weight            2970.424623
      acceleration      15.568090
      model year        76.010050
      origin            1.572864
      dtype: float64
```

```
[24]: # where are the other columns? Check data types
      data.dtypes
```

```
[24]: mpg                float64
      cylinders          int64
      displacement      float64
      horsepower        object
      weight            float64
      acceleration      float64
      model year        int64
      origin            int64
      car name          object
      dtype: object
```

Notice that many columns are of type object, which is not a number. Maybe this has to do with missing values? We know from peeking at the first rows that there are '?' values in there. Let's replace these with the string NaN for not-a-number.

```
[25]: # replace '?' with 'NaN'
      data = data.replace({'?': 'NaN'})
      data.head()
```

```
[25]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	\
0	18.0	8	307.0	130.0	3504.0	12.0	70	
1	15.0	8	350.0	165.0	3693.0	11.5	70	
2	18.0	8	318.0	150.0	3436.0	11.0	70	
3	16.0	8	304.0	150.0	3433.0	12.0	70	
4	17.0	8	302.0	140.0	3449.0	10.5	70	

	origin	car name
0	1	"chevrolet chevelle malibu"
1	1	"buick skylark 320"
2	1	"plymouth satellite"

```

3      1      "amc rebel sst"
4      1      "ford torino"

```

Pandas knows that ‘NaN’ probably means that numbers are missing. Now we can convert the data type from object to float

```

[26]: # convert dtypes
data['horsepower'] = data['horsepower'].astype('float')
data.dtypes

```

```

[26]: mpg          float64
      cylinders    int64
      displacement float64
      horsepower   float64
      weight       float64
      acceleration float64
      model year   int64
      origin       int64
      car name     object
      dtype: object

```

We could have loaded the data with the `na_values` argument to indicate that ‘?’ means missing number:

```

[27]: data = pd.read_fwf('auto-mpg.data', na_values='?',
      ↪names=['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
      ↪'acceleration', 'model year', 'origin', 'car name'])
data.dtypes

```

```

[27]: mpg          float64
      cylinders    int64
      displacement float64
      horsepower   float64
      weight       float64
      acceleration float64
      model year   int64
      origin       int64
      car name     object
      dtype: object

```

This worked nicely. Now we can describe all columns, meaning printing basic statistics. Note that by default Pandas ignores NaN, whereas Numpy does not.

```

[28]: data.describe() # ignores NaN

```

```

[28]:      mpg  cylinders  displacement  horsepower  weight \
count  398.000000  398.000000    398.000000    392.000000  398.000000
mean    23.514573    5.454774    193.425879    104.469388  2970.424623

```


std	7.815984	1.701004	104.269838	38.491160	846.841774
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	model year	origin
count	398.000000	398.000000	398.000000
mean	15.568090	76.010050	1.572864
std	2.757689	3.697627	0.802055
min	8.000000	70.000000	1.000000
25%	13.825000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

We could be interested by these statistics in each of the genders. To get these, we first group values by gender, then ask for the description. We will only look at age for clarity

```
[29]: data.groupby(by='origin').describe().mpg
```

```
[29]:
```

	count	mean	std	min	25%	50%	75%	max
origin								
1	249.0	20.083534	6.402892	9.0	15.0	18.5	24.00	39.0
2	70.0	27.891429	6.723930	16.2	24.0	26.5	30.65	44.3
3	79.0	30.450633	6.090048	18.0	25.7	31.6	34.05	46.6

1.8 Find NaNs

How many NaNs in each column?

We can ask which entries are null, which produces a boolean array

```
[30]: data.isnull()
```

```
[30]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
..	
393	False	False	False	False	False	False	
394	False	False	False	False	False	False	
395	False	False	False	False	False	False	
396	False	False	False	False	False	False	
397	False	False	False	False	False	False	

	model year	origin	car name
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
..
393	False	False	False
394	False	False	False
395	False	False	False
396	False	False	False
397	False	False	False

[398 rows x 9 columns]

Applying `sum()` to this boolean array will count the number of `True` values in each column

```
[31]: data.isnull().sum()
```

```
[31]: mpg                0
      cylinders          0
      displacement      0
      horsepower        6
      weight            0
      acceleration      0
      model year        0
      origin            0
      car name          0
      dtype: int64
```

We get complementary information from `info()`

```
[32]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             398 non-null   float64
1   cylinders        398 non-null   int64
2   displacement     398 non-null   float64
3   horsepower       392 non-null   float64
4   weight           398 non-null   float64
5   acceleration     398 non-null   float64
6   model year       398 non-null   int64
7   origin           398 non-null   int64
```

```

      8   car name      398 non-null   object
dtypes: float64(5), int64(3), object(1)
memory usage: 28.1+ KB

```

We can fill (replace) these missing values, for example with the minimum value in each column

```
[33]: data.fillna(data.min()).describe()
```

```

[33]:          mpg  cylinders  displacement  horsepower  weight \
count  398.000000  398.000000   398.000000   398.000000  398.000000
mean    23.514573    5.454774   193.425879   103.587940  2970.424623
std     7.815984    1.701004   104.269838    38.859575   846.841774
min     9.000000    3.000000    68.000000    46.000000  1613.000000
25%    17.500000    4.000000   104.250000    75.000000  2223.750000
50%    23.000000    4.000000   148.500000    92.000000  2803.500000
75%    29.000000    8.000000   262.000000   125.000000  3608.000000
max    46.600000    8.000000   455.000000   230.000000  5140.000000

          acceleration  model year      origin
count    398.000000   398.000000   398.000000
mean     15.568090    76.010050    1.572864
std       2.757689     3.697627    0.802055
min       8.000000    70.000000    1.000000
25%     13.825000    73.000000    1.000000
50%     15.500000    76.000000    1.000000
75%     17.175000    79.000000    2.000000
max     24.800000    82.000000    3.000000

```

1.9 Count unique values (a histogram)

We finish off, with our good friend the histogram

```
[34]: data['mpg'].value_counts()
```

```

[34]: 13.0    20
      14.0    19
      18.0    17
      15.0    16
      26.0    14
      ..
      31.9     1
      16.9     1
      18.2     1
      22.3     1
      44.0     1
      Name: mpg, Length: 129, dtype: int64

```

```
[ ]:
```

[]: